

Business Model

How this app makes money

A costed, sourced plan for turning Kitto into revenue — what to charge, what to give away, what it costs to serve, and the order to do it in.

01	Where you stand today	Product & code reality
02	The cost floor	Measured AI economics
03	The market	What this category supports
04	The pricing model	Recommended approach
05	Packaging	Free vs paid, mapped to what's built
06	Revenue lines	Ranked by cash speed
07	Cost engineering	Right-sized to zero users
08	Payment rails	Getting paid at all
09	The 90-day plan	Sequence
10	Risks & kill criteria	When to stop

28 July 2026

Written

Measured

Costs, against this repo

Sourced

Market figures, with URLs

Supersedes

Pricing in competitive-todoist.md

The recommendation, in one page

CHARGE

\$8.99/mo

or **\$59/year**. Regional pricing on, so Egypt sees ~\$2–3 automatically.

GIVE AWAY

Everything except AI

A free tier that stores files costs nothing. One that burns Gemini is a hole.

DO FIRST

Sell 10 by hand

\$99 concierge setups. No code. Cash in days, and it tells you who your market is.

Kitto cannot currently charge anyone — not because of a missing price, but because `resolveTier()` returns `'free'` for every user and no payment provider exists. Meanwhile the measured AI cost of a heavy user is **\$9.34/month**, which exceeds what most of the plausible price points return. Both are fixable, and both must be fixed before a single sale.

The five decisions

DECISION	RECOMMENDATION	BECAUSE
What to gate	Gate AI as a capability , not "AI questions per day"	Nobody in a life-admin app approaches 30/day. The current paywall fires only on abusers.
Price	\$8.99/mo · \$59/yr , regional pricing enabled	\$49/yr nets a 5% margin against measured COGS. \$4.99 is loss-making at every usage level.
Free tier	The app without AI, forever + a 7-day full trial	Hard paywall converts 10.7% vs 2.1% freemium, with no retention penalty.
Metering	Flat price + invisible fair-use ceiling . Not credits.	Credits make users ration, and rationing kills the daily loop that retention depends on.
First money	Concierge setups at \$99 , by hand, before writing code	Zero engineering, cash in days, and it answers "who is this for?" — which is still open.

THE NUMBER THAT SHOULD WORRY YOU MOST

Only **17.3%** of subscription apps reach \$1K MRR and **4.6%** reach \$10K MRR within two years. Apps launched 2025 or later hold **3%** of all subscription revenue; pre-2020 apps hold **69%**. Kitto as a consumer subscription is unlikely to replace an income within six months. That is not a judgement on the product — it is the shape of the outcome distribution. This plan is therefore built to find the market cheaply and generate cash early, rather than to optimise a funnel that has no traffic in it yet.

The strategic move is not to price the app better. It is to stop being a consumer subscription business *only* — because the two closest products in this category both ended up distributed by institutions that already owned the customer.

Where you stand today

Three facts from the code, not opinions. Each one blocks revenue.

You cannot sell anything, at any price

SERVER/SRC/MODULES/AI/ROUTES.TS:26

```
// For v1 every user is on the free tier – Pro upgrade comes later.
function resolveTier(): 'free' | 'pro' {
  return 'free'
}
```

Four more sites hardcode `tier: 'free'`. Nothing anywhere sets `subscription.tier = 'pro'`. `me.billing.ts` returns an empty invoice array, and `UpgradeSheet.tsx` honestly says *"Pro isn't open yet."* Even with checkout wired tomorrow, a paying customer would receive free limits. Selling is a build task, not a pricing decision.

There is no way for anyone to find you

`app/page.tsx` is an auth gate that routes to `/welcome` or `/dashboard`. There is no landing page, no marketing surface, no waitlist — no acquisition surface of any kind. Distribution is the binding constraint on this business, and it currently has zero infrastructure behind it.

The product is real, and roughly 80% of a v1

CAPABILITY	STATUS	NOTE
Voice note → matters	BUILT	Background transcription + extraction
Document scan → typed extraction	BUILT	Captures <code>issuer</code> , <code>documentType</code> , per-candidate <code>dueAt</code>
AI chat with tool calling	BUILT	11 tools, user-confirmed state changes
Matters workspace + domains	BUILT	Six life domains, auto-categorised
Local reminders	BUILT	Capacitor local notifications
Account surface + usage meters	BUILT	Reads the live quota endpoints
Daily digest	PARTIAL	Date-bucketed list. No AI ranking, no morning scheduler.
Household sharing	NOT BUILT	No member/invite/household model exists at all
Payments	NOT BUILT	No provider, no entitlements, no tier resolution

THE GAP THAT MATTERS FOR RETENTION

Per `features.md`, the promised **AI-ranked morning digest is not built** — the screen ships as a date-bucketed list, with no ranking model and no scheduler. This is the daily-value loop. AI apps churn roughly 30% faster than non-AI apps, and novelty only sells the first month. **This gap is a bigger threat to the business than any pricing decision in this document.**

One quiet asset

`ScannedDocument` already captures `issuer`, `documentType` and per-candidate `dueAt` — *who your current provider is, and when your policy renews*. That combination is the highest-intent signal in the insurance lead market. Part 06 puts a number on it. Adding structured `amount` / `premium` fields now is a small change that keeps a large option open.

The cost floor

Measured against this repo — prompts extracted and counted at runtime, tool schemas serialised, media billed per Gemini's published rules. Not estimated.

Everything runs on `gemini-2.5-flash` (\$0.30/M input, \$1.00/M audio in, \$2.50/M output). **11 Gemini call sites** exist.

The fixed overhead nobody sees

Every chat round-trip resends this, in full, before any user data:

COMPONENT	CHARS	TOKENS
System prompt (incl. <code>TOOL_RULES</code>)	16,742	4,186
Tool schemas — 11 functions, serialised	15,250	3,813
Prefill few-shot — 58 turns	3,764	941
Floor, before a single word of user data	35,756	8,939

There is no context caching anywhere — `cachedContent` appears zero times. A five-round turn re-bills this floor five times: ~44,700 tokens of pure boilerplate for one exchange.

Cost per operation

OPERATION	ROUND-TRIPS	INPUT TOK	COST	NOTE
Chat — greeting ("hi")	1	8,940	\$0.0031	Tools + prefill still sent
Chat — typical grounded	1	11,563	\$0.0045	20 tasks + 5 voice + history
Chat — one tool call	2	23,463	\$0.0083	
Chat — <code>queryTasks</code> default 50	2	40,270	\$0.0136	Full Mongoose documents
Chat — worst case	5	279,319	\$0.0863	19× typical. Doubles if it ends in a confirmation
Voice note — 60 seconds	2	9,237	\$0.0063	Charges zero quota
Voice note — 5 MiB max	2	37,462	\$0.0358	Byte cap only, no duration check
Doc scan — 20-page PDF	1	6,860	\$0.0041	Max allowed pages
Doc scan — 12MP photo	1	7,892	\$0.0044	24 image tiles, no downscaling
Doc scan — 48MP photo	1	24,404	\$0.0093	88 tiles
Semantic search	1	7,400	\$0.0030	\$0.0163 at the 300-task pool cap
Daily digest prose	1	3,200	\$0.0022	Lazy, but any edit invalidates it
Summarize / categorize / estimate	1–2	900–3,400	\$0.0013–0.0030	Categorize burns 6 calls per click
Reminder planner / clarification	1	160–360	\$0.0002–3	Best-configured sites

What a user costs per month

LIGHT

\$1.10

2 chats + 1 voice note a day

ENGAGED — PLAN FOR THIS

\$3.29

10 chats, 3 voice notes, 2 digests, 1 search daily

HEAVY

\$9.34

Hits the 30/day cap + 10 voice notes

Add 20–60% on top of all three — see the thinking-budget finding below.

The four findings that cost you the most

1 · VOICE NOTES ARE COMPLETELY UNMETERED

`POST /me/voice-notes` has a rate limiter but **never calls** `admitWithinQuota` (`me.voiceNotes.ts:76`) — while the near-identical `/ai/voice/transcribe` does gate. At 12/min sustained, one account can push **17,280 notes/day: \$108/day** at 60-second notes, **\$619/day** at max size, **~\$1,300/day** with uncapped thinking. Per account — and `requireAuth` only verifies the JWT signature, never loading the user, so `emailVerifiedAt` is written by the verification route and **read by nothing**. Throwaway emails multiply it without limit.

2 · THINKING BUDGET MISSING ON 8 OF 11 CALL SITES

Only three set `thinkingConfig`. The other eight run dynamic thinking billed at the **\$2.50/M output rate** — and the gating is exactly inverted: the two sites that *do* disable thinking have the two smallest prompts (634 and 737 chars), while the largest prompt in the codebase (`voiceCore/extract.ts`) has none. At the 24,576-token ceiling that is **\$0.061 of pure surcharge per call**. `maxOutputTokens` appears **zero times** in the entire codebase.

3 · ONE QUOTA UNIT CAN BUY 24 BILLED CALLS

`withGeminiRetry` allows 3 attempts and **429 is a retry trigger** — a rate-limited call gets retried twice more, each billed. Job-level `maxAttempts: 4` nests on top. Separately, `refineWithAi` fires per reminder task with no quota, no rate limit and no attribution: one chat turn creating 5 reminders silently fires 5 extra Gemini calls.

4 · YOUR TELEMETRY UNDER-REPORTS BY 4.2×

`service.ts:367` assigns `usage = ev.usage` inside the round loop, so each round overwrites the last. A 5-round turn reports ~66,854 tokens against 279,319 actual. Whatever dashboard you build on this will lie to you until it is fixed.

AND ONE PIECE OF GOOD NEWS

`hallucinationGuard.ts` is 26 lines of pure regex — no second model call, zero cost. The daily digest is lazy rather than a cron, so dormant users cost nothing. And `admitWithinQuota` is already a correct, TOCTOU-safe atomic admission — the hard part of metering is **already written**, just pointed at the wrong limits.

The market

What this category actually supports, and where its ceiling is.

Three price bands, and Kitto straddles two

BAND	PRICE	WHO IS THERE
Consumer to-do	\$3–10/mo	Todoist \$7 · Any.do \$7.99 · TickTick \$3.99 · Tiimo ~\$10 · Structured \$6.99
Family / household admin	\$39–120/yr <i>flat</i>	Cozi \$39/yr <i>whole household</i> · Maple \$40 · Everplans \$99.99 · Prisdio \$120
AI "does work for you"	\$17–49/mo	Sunsama \$17–22 · Akiflow \$19–34 · Motion \$19–34 · Martin \$21–49
Document vault ceiling	\$40/mo · \$480/yr	Trustworthy Platinum — <i>and it includes 3 hours of human concierge</i>

The closest analogs — and two corpses

PRODUCT	PRICE	WHAT IT TELLS YOU
Trustworthy	Free 2GB / \$10 / \$20 / \$40 mo	The key comp. AI is the mid-tier upsell , not the base product.
Everplans	\$99.99/yr — free tier is 3 items	Sells at \$27/yr via VSP benefits . Acquired twice. Drifted hard to B2B2C.
Cozi	\$39/yr, whole household	Household-flat pricing works and is well understood.
Todoist + Ramble	\$7/mo, voice unlimited	Voice→task is now a \$5–7 commodity, with Google subsidising the inference.
Yohana (Panasonic)	\$249/mo — dead 30 Sep 2025	Human concierge at the category ceiling. Still died.
Milo	Dead	Founder: the technology was <i>"too early to be reliable in any useful way."</i>
Warranty trackers	~\$0.99/yr for ad removal	A graveyard. No standalone consumer willingness-to-pay . Feature, never product.

The benchmarks that should drive your decisions

RevenueCat *State of Subscription Apps 2026* — 115,000+ apps, \$16B+ revenue.

METRIC	NUMBER	WHAT IT MEANS FOR YOU
Hard paywall vs freemium	10.7% vs 2.1%	5× conversion, 8× revenue per install, no retention penalty (27% vs 28% at 1yr)
Download → paid, median	2.0%	NA 2.8% · India/SEA 0.7% — region changes everything
High-priced vs low-priced conversion	2.8% vs 1.4%	Year-1 LTV \$62.19 vs \$10.69 . Cheap does <i>not</i> convert better.
Trial → paid by length	25.5% → 42.5%	≤4 days 25.5% · 5–9 days 37.4% · 17–32 days 42.5%. Longer wins.
Productivity: converts on Day 0	71.9%	The decision is instant. Anything needing arithmetic at that moment costs sales.
AI apps: revenue vs churn	+41% / –30%	\$30.16 vs \$21.37 per payer, but annual retention 21.1% vs 30.7%
Annual Year-1 churn	~72%	And 35% of annual cancellations happen in Month 1
Median pricing	\$10.00 mo · \$34.80 yr	Annual discounts cluster 25–45%
Reach \$1K MRR / \$10K MRR in 2 yrs	17.3% / 4.6%	Plan for a long grind

The ceiling of this category is \$40/month, and you cannot reach it without humans in the loop. Two companies proved that — one by charging \$249/month for concierge and dying anyway, the other by trying to replace those humans with AI.

The pricing model

The recommended approach, and the three alternatives that were considered and rejected.

The structure

TIER	WHAT IT IS	PRICE
Free, forever	The app without AI — manual matters, reminders, document storage (~20 docs), full export	\$0
Trial	Full AI, everything unlocked	7 days free
Kitto Plus	The AI — scan→extract, voice→matters, chat, digest, unlimited documents	\$8.99/mo · \$59/yr
Top-ups <i>(optional)</i>	Only when fair use is exceeded	+20 scans \$4.99 · bulk import 200 docs \$29

Why this shape

The free tier costs you nothing to serve

This is the whole trick. A free tier that burns Gemini is a hole; one that only stores files is a rounding error. Users keep their data forever — which is the trust promise a document vault lives on — and they upgrade when they want the magic, not when they hit an arbitrary counter.

Gate AI as a capability, not as a counter

Your current model gates **30 AI messages/day free vs 300 Pro**. Both numbers are wrong in the same direction: no real life-admin user approaches 900 messages a month, so the paywall **never fires for genuine users — only for abusers** — and 300/day is a ceiling nobody reaches, so Pro visibly sells nothing. You are gating a *cost input* where you should be gating a *capability*.

\$59/year is the headline; \$8.99/month is the anchor

They are the same product and they feel completely different. The 45% annual discount sits inside the normal 25–45% band, and annual front-loads cash — which matters when you have none.

WHY \$59 AND NOT \$49

\$49/yr nets **\$3.47/mo** after a 15% store cut, against a measured **\$3.29/mo** engaged-user cost — a **5% margin**, and still negative on heavy users even after the Part 07 fixes. The category median pulls toward \$49; your COGS does not permit it. \$59 nets \$4.18/mo. Selling via the web through a merchant-of-record (~5% rather than 15%) adds back roughly \$0.45/mo.

Turn on store regional pricing — it resolves the question you can't answer yet

Apple and Google adjust by market automatically: **\$8.99 in the US displays as roughly \$2–3 in Egypt**. NA converts at 2.8% and India/SEA at 0.7%, so this is how you serve both without wrecking either — and it means you do not

have to decide your target market before launch.

Margin at each price

PRICE	YOU KEEP	ENGAGED TODAY	ENGAGED POST-FIX	HEAVY TODAY	HEAVY POST-FIX
\$4.99/mo	\$4.24	\$0.95 (22%)	\$1.94	-\$5.10	-\$1.76
\$59/yr (= \$4.18/mo net)	\$4.18	\$0.89 (21%)	\$1.88	-\$5.16	-\$1.82
\$8.99/mo	\$7.64	\$4.35 (57%)	\$5.34	-\$1.70	+\$1.94
\$9.99/mo	\$8.49	\$5.20 (61%)	\$6.19	-\$0.85	+\$2.49

15% store cut assumed. "Today" figures are measured; "post-fix" applies the measured 20–60% thinking-budget saving conservatively (chat already has thinking configured).

READ THE ANNUAL ROW CAREFULLY

A heavy user on the \$59 annual plan is loss-making even after the fixes. This is exactly what the fair-use ceiling exists to stop — without it, annual pricing is an open invitation to your most expensive users. Annual subscribers do skew lighter on average (buy and forget), so the blended figure is better than this row implies. But the tail is real, and it must be capped rather than hoped away.

Considered and rejected

Credits as the headline model **REJECTED**

Credits make COGS track revenue — the property a founder who cannot absorb loss most wants — and they would permit a lower entry price. But every consumer product using them well is B2B or prosumer: Motion at \$34/mo (7,500 credits/seat, \$0.25/100 overage), Notion at \$10/1,000 credits — both expensed by someone other than the user. At the \$5–10 consumer tier, flat pricing dominates, for three reasons specific to Kitto:

- **Credits cause rationing.** Less usage → weaker habit → more churn. AI apps already churn ~30% faster. Suppressing the loop is the wrong trade.
- **"500 credits" is unanswerable at the point of sale.** Productivity converts **71.9% on Day 0** — anything requiring arithmetic at that instant costs conversions.
- **It contradicts the product test.** `principles.md` asks: *does this make life admin feel less stressful, not more?* A balance ticking down while you scan an insurance policy is the precise opposite of the relief this product sells.

ADOPTED INSTEAD — THE INVISIBLE CEILING

Flat price, with fair use set at **5–10× normal usage** (≈40 scans, 100 voice notes, 300 chat messages per month) so ~95% of users never discover it exists. On breach, offer a top-up rather than a wall. This buys the cost ceiling without the credit anxiety — the pattern Otter, Voicenotes and Granola all use. **And admitWithinQuota** already implements it correctly; you are changing numbers, not building a metering system.

Where credits genuinely are right: bulk import ("digitize my filing cabinet" is ~200 documents and a real \$10–20 COGS event — sell it as a \$29 pack), and B2B2C seats if the broker route opens.

A lifetime deal **REJECTED**

Tempting for fast cash, and genuinely dangerous. **ChatPlayground AI revoked every lifetime licence it had sold and demanded \$875 from existing lifetime owners**, citing API costs. Where lifetime survives in AI products it is a credit bundle or bring-your-own-key, never unlimited. Structured can sell \$99.99 lifetime because it has no AI COGS. You cannot. **Prepaid two-year terms are the survivable version** if you need cash up front.

Pricing at \$4.99 to feel accessible **REJECTED**

High-priced apps converted **2.8% vs 1.4%** and returned **\$62.19 vs \$10.69** Year-1 LTV. Cheap does not convert better — and per the margin table, \$4.99 is loss-making on heavy users even after every cost fix in this document. If the price still feels high, the answer is **regional pricing and the \$59 annual framing**, not a lower global number.

Packaging

Free vs paid, mapped strictly to what exists in the codebase today.

CAPABILITY	BUILT?	FREE	PLUS
Manual matters, domains, reminders	YES	✓	✓
Document storage + vault	YES	20 documents	Unlimited
Data export (JSON)	YES	✓ always	✓
Document scan → extraction	YES	—	40/mo fair use
Voice note → matters	YES	—	100/mo fair use
AI chat + tool calling	YES	—	300/mo fair use
Daily digest	PARTIAL	—	✓
Auto-categorise	YES	—	✓
Clarifications / uncertainties	YES	—	✓
Household sharing	NO	A real build, not a repackaging. Park it — see Part 09.	

KEEP EXPORT FREE, PERMANENTLY

It is the cheapest trust signal you have. Gating data export in a product whose entire promise is *"your important documents are safe here"* undercuts the exact thing you are selling. Trustworthy and Everplans both compete on trust; neither holds data hostage.

The upgrade moment

The paywall should fire at the moment of demonstrated value, not at a counter. For Kitto that is: the user scans their first document and watches the renewal date, issuer and amount populate themselves. That is the *"oh"* moment. The trial exists so that moment happens before the ask — and the ask should reference what just happened, not a feature list.

Revenue lines

Ranked by cash speed for someone with no runway — not by size of the opportunity.

1 Concierge setup — do this first

"Send me your documents; I'll set up your vault and your renewal calendar. \$99."

Requires zero code. Cash in days. And it answers the question that is still open: *who actually has this pain badly enough to pay for it?* Ten of these is \$990 and a real customer-discovery dataset. This is the highest-value line in the list for your specific situation, and it is the one most founders skip because it does not feel like building a product.

2 Annual subscriptions from one narrow launch

50 payers × \$59 = **\$2,950**. Requires checkout, entitlements and a landing page — none of which exist today. This is Part 09, weeks 3–12.

3 B2B2C — the pattern both close analogs converged on

Everplans sells white-labelled through advisors and insurance agents at **\$27/yr via VSP** versus \$99.99 direct, and has been acquired twice. Prisdio runs through AARP. Wealth.com's vault is advisor-only. **Consumer document vaults consistently struggle to acquire directly and end up distributed by institutions that already own the customer.**

One broker with 300 clients at \$2/client/month is **\$600/month from a single sale**. It needs multi-tenancy you have not built — but it needs a *demo*, not a finished app, and the sales cycle runs in parallel with everything else.

4 Insurance renewal referrals — biggest ceiling, phase 3

LEAD TYPE	US PRICE	NOTE
Shared auto / home	\$10–25 / \$7.50–30	~25% contact rate
Exclusive auto / home	\$25–50	~70% contact rate
Live transfer, auto/home	\$30–75	Life/health \$50–120+

Your entire \$59 annual subscription is worth 2–6 referrals, and a household with car + home + health policies generates four renewal events a year. MediaAlpha did **\$310M in Q1 2026** on ~11.7M monthly referrals; EverQuote \$190.9M. And you already capture the signal: `issuer` + `dueAt` is *who your provider is and when you renew* — the highest-intent data in that market.

THREE CONSTRAINTS, ALL SERIOUS

Legal (US): TCPA governs consent, and one-to-one consent rules make blanket data-sharing flows hazardous.

Legal (UK): introducing insurance business is an FCA-regulated activity requiring authorisation or appointed-representative status. **Trust:** selling renewal data derived from documents users handed you for safekeeping is a cliff, not a slope — Trustworthy and Everplans both notably decline to monetise this way. If you ever do it, it must be an explicit, user-initiated *"get me a quote"* button, never background monetisation.

Cheap option to keep now: add structured `amount / premium` fields to the extraction schema. Small change today; the difference between having this option later and not.

Cost engineering

The audit found 13 issues. Most are scale problems. At zero users, context caching saves nothing — do these, in this order, and stop.

#	FIX	EFFORT	WHY
0	Hard billing cap on the Gemini / GCP project	10 min, no code	Bounds your maximum possible loss regardless of everything below
1	Quota on <code>me.voiceNotes.ts</code>	~3 lines	Closes the \$108–1,300/day/account hole
2	Wire <code>resolveTier()</code> to the real subscription	Small	Without this, Plus is unsellable at any price
3	<code>thinkingBudget: 0</code> on the 8 sites missing it	~8 one-line edits	20–60% off everything outside chat
4	Downscale images to 1024px long edge	Small	12× fewer media tokens (6,192 → 516)
5	<code>maxOutputTokens</code> — appears zero times today	Small	Bounds the tail

Only three are strictly required to launch: 0, 1 and 2. Items 3–5 are what move you from break-even to a real margin, and they should land before the first annual subscriber, not before the first line of marketing.

THE TRADE YOU ARE MAKING

Price and cost work are linked — you can defer one only by accepting the other. **\$4.99 only works if you do all the cost work, and even then it doesn't clear heavy users. \$8.99 works even if you skip most of it.** That is the real argument for the higher price: not market positioning, arithmetic.

Deferred until you have users

- Context caching — saves nothing at zero traffic, meaningful at scale
- Lean-projecting `queryTasks` — 50 rows currently serialises 13,088 tokens of full Mongoose documents, including `userId` on every row
- The `usage` overwrite at `service.ts:367` (under-reports 4.2×)
- Retry-storm capping, and categorize's uncapped tag vocabulary

ONE INFRASTRUCTURE ITEM THAT IS NOT OPTIONAL

Documents are written to **local disk** (`documentScanStorage.ts:23` → `uploads/document-scans`), which is ephemeral on most container hosts. **Taking money for "your documents are safe here" while storing them on a container filesystem is the one risk in this document that could generate actual liability.** Move to object storage (Cloudflare R2 has no egress fees) before the first payment — not before the first user.

Payment rails

Being able to receive the money is a prerequisite, and it is not straightforward.

STRIPE DOES NOT OPERATE IN EGYPT

Confirmed against Stripe's own country list. Egypt appears neither in the 60+ supported countries nor in the Paystack extended network. This is definitive, and it rules out the default path most tutorials assume.

RAIL	EGYPT	FEE	NOTE
Paddle (merchant of record)	Listed	5% + 50¢ all-in, tax handled	Verify seller eligibility directly before building on it
Dodo Payments (MoR)	Explicit Egypt page	4% + 40¢ (+1.5% intl, +0.5% subs)	Markets to developing-market founders
Lemon Squeezy	—	—	Do not build on it — Stripe-acquired, migrating to Stripe Managed Payments
Stripe direct	Not available	—	Confirmed

Take-home by channel, solo dev under \$1M/yr

CHANNEL	YOU KEEP	NOTE
Web via merchant-of-record	93–96%	The best channel — push annual plans here
iOS / Android in-app	~84%	15% Small Business Program + 1% RevenueCat
iOS EU	80–90%	10–15% + 5% Core Technology Commission
iOS US via external link	~97% <i>today</i>	Live litigation — see below

DO NOT BUILD A MODEL THAT DEPENDS ON THE EXTERNAL-LINK RATE

Apple currently collects **0%** on US external-link purchases, but this is unsettled law, not policy. Timeline: Apr 2025 contempt ruling barred any commission → Dec 2025 Ninth Circuit affirmed contempt but *vacated the total ban* → **Jun 2026 the Supreme Court granted certiorari**, with a decision expected late 2026 or early 2027. Today's 0% holds. Treat it as temporary upside, never as the basis of your unit economics.

THE US LLC ROUTE CARRIES A \$25,000 TRAP

Stripe Atlas is \$500 one-time plus ~\$100/yr, and it works. But a foreign-owned US LLC must file **Form 5472** plus a **pro forma 1120 annually, even at zero revenue**. The penalty for missing it is **\$25,000 per form per year**, plus \$25,000 per additional 30-day period after IRS notice. Do not take this on until web revenue justifies the compliance burden.

Recommended stack: Paddle or Dodo as merchant-of-record for web → RevenueCat for mobile entitlements (free under \$2,500 monthly tracked revenue, then 1%) → store billing at 15% under \$1M. Add the US LLC only when web revenue justifies it.

The 90-day plan

Sequenced so that you learn before you build, and get paid before you scale.

Weeks 1–2 • Stop the bleeding, learn something

- ☐ **Set the billing cap.** Ten minutes. Pick a number you can stand losing.
- ☐ **Add the voice-note quota.** Three lines, calling a function you already wrote.
- ☐ **Run 5–10 concierge setups at \$99, by hand.** The goal is not the money. It is finding out which market flinches at the price and which says *finally*.

Weeks 3–6 • Make it sellable

- ☐ Wire `resolveTier()` to the real subscription — nothing else matters until this works
- ☐ Add checkout (Paddle/Dodo web first, store billing second)
- ☐ Move document storage off local disk to object storage
- ☐ `thinkingBudget: 0` and image downscaling
- ☐ **Build a landing page** — you currently have no acquisition surface whatsoever

Weeks 7–12 • One narrow launch

- ☐ Pick the single market that responded best in weeks 1–2
- ☐ Launch at \$8.99 / \$59 with regional pricing on and a 7-day trial
- ☐ Open one B2B2C conversation in parallel — it needs a demo, not a finished product
- ☐ **Target: 50 paying users.** Not 5,000.

WHY CONCIERGE COMES BEFORE CODE

You do not currently know who your first paying users are. Every engineering decision downstream of that — which market's regional price to optimise, whether to build household sharing, whether the expat wedge is real — is a guess until you do. Ten conversations at \$99 each buys you that answer *and* pays for the month.

Risks & kill criteria

Decide these now, while you are not emotionally invested in the answer.

SIGNAL	WHAT IT MEANS
10 concierge offers, zero takers at \$99	Nobody has this pain at a price that sustains you. That is a market signal, not an execution failure.
Landing page live 60 days, <100 signups	The positioning is not landing. Change the wedge before touching the product.
50+ trials, <5% convert	The product does not deliver its promise yet. More features will not fix it — the hero loop is not closing.
Month-1 churn above 50% on annual	They bought a promise the product did not keep. Baseline is already rough: 72% Year-1 annual churn, and 35% of annual cancellations land in Month 1.

The risk to watch hardest

AI apps earn **41% more per payer** (\$30.16 vs \$21.37 Year-1) but churn far faster — **21.1% vs 30.7%** annual retention. Novelty sells the first month; only the daily value loop keeps them. For Kitto that loop is the morning digest and the renewal warnings — and **the AI-ranked digest is not built**. The screen ships as a date-bucketed list, with no ranking model and no scheduler.

That gap is the single biggest threat to this business, and it is more important than every pricing decision in this document combined. A correctly-priced product that people stop opening in week three is still a dead product.

Other standing risks

- Trust is one wrong date away. `principles.md` already names this — one bad renewal reminder loses a user permanently. The citation primitive is the mitigation, and it must not be traded away for speed.
- Todoist has already commoditised voice capture at \$7/mo with Google subsidising inference. Documents, domain depth and execution are the defensible ground; voice capture alone is not.
- Warranty tracking has no standalone willingness-to-pay. Keep it as a feature, never position on it.
- Concentration risk on Gemini. One provider, one API key, no fallback path. A price change or policy change is an unhedged event.

Sources. Market: RevenueCat *State of Subscription Apps 2026* (115k+ apps) · Trustworthy, Everplans, Cozi, Todoist, Any.do public pricing pages · Stripe country list · Maverick Marketing insurance lead costs (May 2026) · MediaAlpha and EverQuote Q1 2026 results · Ninth Circuit *Epic v. Apple* and June 2026 cert grant.

Costs were measured directly against this repository — prompts extracted and counted at runtime, tool schemas serialised, media billed per Gemini's published rules (32 tokens per second of audio, 258 tokens per 768x768 image tile). Token counts use the chars÷4 heuristic. `gemini-2.5-flash` list pricing: \$0.30/M input, \$1.00/M audio in, \$2.50/M output.

Flagged as unverified: Paddle Egypt seller eligibility (highest-stakes open item — confirm with Paddle directly) · Apple payouts to Egyptian bank accounts · Tiimo exact pricing (vendor page unreachable) · UK/EU/MENA insurance lead values (US data only) · post-fix cost figures are derived estimates, not measurements.